

0 Index

0	Index	1
1	Overview	4
2	Introduction	6
2.1	Situation overview	6
2.2	Purpose	7
2.3	Scope	8
2.4	Project goals	8
3	Theoretical aspects	10
3.1	Assignment problem	10
3.2	Greedy algorithms	14
3.3	Heuristics and metaheuristics	18
3.3.1	Evolutionary Computation	18
3.3.2	Evolution Strategies	20
3.3.3	Genetic algorithms	22
3.3.3.1	Selection	23
4	Problem definition	28
5	Proposed solution	31
5.1	Search space	31
5.1.1	Assignments	31
5.1.2	Solutions	31
5.1.3	States	32
5.1.4	Instances	33
5.2	Collisions	33
5.2.1	Lazy Collision Matrix	34
5.3	Classroom filters	36
5.3.1	Lazy Filter Dictionary	37
5.4	Greedy algorithm	39
5.4.1	Preprocessing	39
5.4.2	Heuristic	39
5.4.3	Repairs	39
5.5	Genetic algorithm	39
5.5.1	Fitness function	39
5.5.2	Operators	39
5.5.2.1	Selection	39
5.5.2.2	Crossover	39
5.5.2.3	Mutation	39
5.5.2.4	Tournament	39
5.5.3	Parameters	39

6	Project planning and budget overview	40
6.1	Planning	40
6.2	Budget summary	40
7	Analysis	41
7.1	System definition	41
7.2	System requirements	41
7.2.1	Interface	41
7.2.2	Input	41
7.2.3	Configuration	42
7.2.4	Algorithm	42
7.3	Subsystem mapping	43
7.4	Preliminary class diagram	43
7.5	Analysis of use cases	43
7.6	Analysis of user interfaces	43
7.7	Test plan specification	43
8	System design	44
8.1	System architecture	44
8.2	Class design	44
8.3	Interaction and state diagrams	44
8.4	Activity diagram	44
8.5	Interface design	44
8.6	Technical specification of the test plan	44
9	System implementation	45
9.1	Standards and references	45
9.1.1	Standards	45
9.1.2	Licenses	45
9.1.3	Other references	45
9.2	Programming languages	45
9.3	Tools and programs used in development	46
9.4	System development	46
10	Test development	47
10.1	Unit tests	47
10.2	Integration and system tests	47
10.3	Usability and accessibility tests	47
10.4	Performance tests	47
11	Experimental results	48
12	System manuals	49
12.1	Installation manual	49
12.2	Execution manual	49
12.3	User manual	49

12.4 Programmer manual	49
13 Conclusions and future work	50
13.1 Final conclusions	50
13.2 Future work	50
14 Budget	51
14.1 Internal budget	51
14.2 Client budget	51
15 Annexes	52
15.1 Definitions and abbreviations	52
15.2 Submission contents	53
16 Source code	54

1 Overview

This document presents all the important information regarding the *Classroom management at the School of Computer Engineering using Artificial Intelligence methods* end-of-degree thesis.

It is important to note that the structure of the contents for this document is done following the criteria and recommendations of the template document for Degree's and Master's Thesis of the School of Computing Engineering of Oviedo (version 1.4) by Redondo [Red]. However, some additional chapters were introduced in order to capture the particularity of the work carried out, inspired by the research of de la Cruz [dLC18].

Introduction. Here we explain in a simple way the problem we want to solve, what reasons are behind the development of the project and give a description of the current situation of the School with regard to this and other similar problems. The scope and goals of the project are also discussed.

Theoretical aspects. The first chapter delving into the theory supporting the developed system. One example of an assignment problem is presented and solved using a greedy algorithm and a genetic algorithm, which helps to better internalize the concepts.

Problem definition. Here the formulation of the problem as an assignment problem is elaborated.

Proposed solution. This chapter lays down in detail the proposed solution to the previously defined assignment problem by means of a genetic and greedy algorithms.

Project planning and budget overview. For the planning, the Gantt chart of the project is shown, as well as the work breakdown structure. The internal and client budgets are presented, but not how they were calculated. That goes in the budget section.

Analysis. An analysis of the system, with the system requirements, draft diagrams, use cases and the test plan specification.

System design. The technical details of the system analysed in the previous section. Here we include the finalised diagrams, as well as the architecture of the system and the in-depth test plan.

System implementation. Details of the development of the software. The programming languages, standards and tools used to code the system, and all the relevant information gathered in the process of creating the system.

Test development. A rundown of all the testing done for the system, with explanations for every test and the obtained results.

Experimental results. The conclusions reached after experimenting with different input data and the optimal configuration for the genetic algorithm's parameters.

System manuals. All the manuals for the system, with screenshots and guided steps meant to help the target audience for each document.

Conclusions and future work. The conclusions after the implementation of the project are given, as well as a list of possible improvements and new functionalities for the prototype.

Budget. Here we give the full details for the elaboration of the internal and client budgets, with all the intermediate steps that led us to that result.

Annexes. The glossary of definitions and abbreviations and a small commentary on the submitted files.

Source code. The *javadoc* of the software developed.

2 Introduction

The School of Computing Engineering of the University of Oviedo has more than twenty classrooms, including theory classrooms and laboratory classrooms. Each semester there are over three hundred groups, each with their type (theory, seminar or laboratory), subject and schedule. The timetable of the groups varies on a weekly basis, this means that not all groups have to attend classes all weeks, and some of them do not even have repeating patterns.

This makes assigning classrooms to groups a complicated task, since there can be no temporal collisions. When various other constraints enter the equation, such as minimising the number of labs used by a subject or assigning classrooms to Spanish groups that are different from English groups, things become much more complex.

All this assignments are done *manually* by one person. The number of enrolled students can only be *guessed* when this process is done. This means that groups can be created, modified or removed once the semester has already started, so more assignments are usually made, checking once again all the restrictions. These new assignments are difficult to manage as there is not much room for flexibility to change those made before the semester. This is due to the fact that both students and teachers already use the initial assignments as a reference.

This project provides the supervisor of this process with a tool to help them calculating the assignments, reducing their workload. Not only does it generate assignments for all the groups of the semester, but can use previous assignments, total or partial, to calculate a subset of assignments (for example, the assignments for the new groups created in the middle of the semester). On top of that, the prototype developed in this thesis makes finding a set of free classrooms to hold events easy and fast, using the assignments generated previously by the system itself.

2.1 Situation overview

At the beginning of each semester, the School opens a process in which the person in charge takes the list of groups for the semester, their schedules and the list of classrooms, and performs a manual compilation of all the assignments.

There are a number of other similar procedures, like the creation of the exam timetable or the assignments of enrolled students to subject groups. However, some are not manual, but automated by a system, like the previously mentioned procedure of assigning students to groups. Seeing the potential of such tools, I was given the task of automating the assignment of classrooms to subject groups by similar means.

The procedure of assigning the classrooms is done after configuring the student groups for the semester and knowing their schedules. Even though it is a manual process, the supervisor does not start making the assignments from scratch. First, they have the knowledge of previous years, and then they have a list of preferences or premade assignments. For example, certain laboratories can only be assigned to specific groups, like the ones from the Electronic Technology of Computers subject. The system described in this document preserves these sources of information and builds on top of them.

2.2 Purpose

This project aims to help the personnel of the School manage their classrooms. It will address two main functionalities, the automation of the process of assigning classrooms to all the groups of a given semester (starting from scratch or using a previous partial or total assignment), and a tool that searches for gaps in a previous set of assignments for single or multi-day events in one or more classes.

The implementation of this system is intended to assist in the work of the supervisor for this process, and provide an efficient and flexible tool that expands the possibilities of such work. To do so, the program executes two algorithms, a genetic algorithm guided by a greedy algorithm. For a more detailed view on these algorithms the reader might refer to [3](#). Once the assignments have been calculated, the system will allow the users to find classrooms to hold specific events in the middle of the semester.

Along with the system, the system manuals are submitted. These have the purpose of teaching how to install, use, maintain and extend the system. Apart from the manuals, another tool to generate the necessary files for the program is handed.

2.3 Scope

The project needs to formally define the problem of assigning classrooms of the School to all the groups of the semester, conduct a study on the problem and propose a solution.

A development of a software prototype that solves the problem is planned, designed, implemented and tested. This prototype will solve the two main functionalities indicated in 2.2 and will consist of a command line application that takes input data in plain text files and outputs the solution to plain text files. The program is configured by different configuration files depending on the functionality being executed. An experimental study on the results of the software system is carried out, finding the most fitting default values for the configuration files. The project also contains the system manuals of the application, which consist of the installation, usage, user and programmer manuals.

Finally, an additional tool for automating the creation of the input files of the software is given. It uses a format agreed with the client and will use the same technical specifications of the main prototype, like the programming language and the development environment.

2.4 Project goals

We can identify from the scope the following objectives. They need to be met in order to close the project successfully:

1. Formally define the problem of assigning classrooms to the groups of the School.
2. Study the problem and the means to solve it.
3. Define the proposed solution.
4. Build a prototype that solves the problem using the algorithms described in the proposed solution.
 - (a) It will receive plain text input files with the required data.
 - (b) It will output the solution to plain text files.

- (c) It will be able to make the assignments starting from scratch or from a total or partial set of assignments.
 - (d) It will be able to search a set of free classrooms for a specific event in one or more days.
5. Make a set of experiments to find the best default values for the configuration files.
 6. Write a set of manuals to cover the essentials of the system.
 7. Create another software tool that will automate the creation of the input plain text files for the main system.
 8. Validate solution with the users.

3 Theoretical aspects

A digital magazine Bootaku works with three freelancers. Dante, Virgil and Beatrice. Together they write a section about book reviews. Gathering data from previous sections, Bootaku wants to define and solve a problem of efficiently assign all reviews to the three critics so that the section gets the highest profit. For the assignments, Bootaku wants every book review to have one (and only one) associated freelancer. If a freelancer ends up with no reviews, the assignments are still valid if and only if the previous condition is met.

3.1 Assignment problem

The problem described before is an example of an assignment problem. It can be generalised with the following elements:

A set of n freelancers f

A set of m book reviews r

An assignment matrix of $n \times m$ assignments a_{fr} such that $a_{fr} = 0$ when freelancer f is not assigned to book review r and $a_{fr} = 1$ when freelancer f is assigned to book review r .

A profit matrix of $n \times m$ profits p_{fr} which indicate the profit obtained when assigning freelancer f to book review r and that $p_{fr} > 0$.

A valid solution is defined as a matrix of assignments where all the book reviews have a freelancer assigned to them and no book review has more than one associated freelancer.

The profit for all the assignments will then be:

$$\sum_{f=1}^n \sum_{r=1}^m a_{fr} p_{fr} \quad (3.1)$$

The optimal solution consists on having a set of assignments such that the sum of all the profits for the current assignments is maximised.

For example, imagine that for the next month's section, we have the following data. The information is represented by means of two sets: F for the freelancers and R for the reviews.

$$F = \{Dante, Virgil, Beatrice\} \quad (3.2)$$

$$R = \{Divina Commedia, El Quijote, Voyage au bout de la nuit, Todo modo\} \quad (3.3)$$

Then, our assignments and profits will be represented by the A and P matrices.

$$A = \begin{matrix} & \begin{matrix} DC & EQ & VN & TM \end{matrix} \\ \begin{matrix} Dante \\ Virgil \\ Beatrice \end{matrix} & \begin{pmatrix} a_{11} & a_{12} & a_{13} & a_{14} \\ a_{21} & a_{22} & a_{23} & a_{24} \\ a_{31} & a_{32} & a_{33} & a_{34} \end{pmatrix} \end{matrix} \quad (3.4)$$

$$P = \begin{matrix} & \begin{matrix} DC & EQ & VN & TM \end{matrix} \\ \begin{matrix} Dante \\ Virgil \\ Beatrice \end{matrix} & \begin{pmatrix} p_{11} & p_{12} & p_{13} & p_{14} \\ p_{21} & p_{22} & p_{23} & p_{24} \\ p_{31} & p_{32} & p_{33} & p_{34} \end{pmatrix} \end{matrix} \quad (3.5)$$

Where each row represents a freelancer and each column represents a book review. So freelancer 1 is Dante, 2 is Virgil and 3 is Beatrice. The same goes for the book reviews. Book review 1 is *Divina Commedia*, 2 is *El Quijote*, 3 is *Voyage Au Bout De La Nuit* and 4 is *Todo Modo*.

Now, we are going to study valid and non-valid solutions. As we explained before, a solution is valid if every book review has a freelancer assigned to it, and no more than one.

We will analyse four sets of values for the A matrix:

$$A1 = \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 1 & 0 & 1 \end{bmatrix} \quad (3.6)$$

$$A2 = \begin{bmatrix} 0 & 0 & 0 & 0 \\ 1 & 0 & 1 & 0 \\ 0 & 1 & 0 & 1 \end{bmatrix} \quad (3.7)$$

$$A3 = \begin{bmatrix} 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 \\ 1 & 0 & 0 & 0 \end{bmatrix} \quad (3.8)$$

$$A4 = \begin{bmatrix} 0 & 1 & 1 & 0 \\ 0 & 0 & 0 & 1 \\ 1 & 1 & 0 & 0 \end{bmatrix} \quad (3.9)$$

From these matrices, we can deduce that $A1$ and $A2$ are valid solutions, because they have one freelancer for each book review. We are not concerned with a freelancer having no book reviews assigned. However, a book review without an associated freelancer represents a non-valid solution. That is precisely the case for $A3$, the book review for *El Quijote* has not an assigned freelancer. In the case of $A4$, the fact that *El Quijote* has two freelancers assigned makes it a non-valid solution.

Now, we will give values to the P matrix in order to discuss possible optimal solutions. We will compare them with the assignment matrices $A1$ and $A2$

$$P1 = \begin{bmatrix} 9 & 1 & 5 & 4 \\ 2 & 8 & 14 & 2 \\ 7 & 11 & 10 & 6 \end{bmatrix} \quad (3.10)$$

$$P2 = \begin{bmatrix} 9 & 1 & 5 & 4 \\ 13 & 8 & 14 & 2 \\ 7 & 11 & 10 & 6 \end{bmatrix} \quad (3.11)$$

$P1$ and $A1$:

$$Profit = \sum_{f=1}^n \sum_{r=1}^m a_{fr} p_{fr} = 9 + 11 + 14 + 6 = 40 \quad (3.12)$$

$P1$ and $A2$:

$$Profit = \sum_{f=1}^n \sum_{r=1}^m a_{fr} p_{fr} = 2 + 11 + 14 + 6 = 33 \quad (3.13)$$

We can observe that for $P1$, the assignments defined in $A1$ are better than those in $A2$, because they result in a better profit. Another important remark about $A1$ is that it is the optimal solution to the problem, because it assigns the book reviews to the freelancers with the best profit value for their assigned books. Now $P2$ will be evaluated.

$P2$ and $A1$:

$$Profit = \sum_{f=1}^n \sum_{r=1}^m a_{fr} p_{fr} = 9 + 11 + 14 + 6 = 40 \quad (3.14)$$

$P2$ and $A2$:

$$Profit = \sum_{f=1}^n \sum_{r=1}^m a_{fr} p_{fr} = 13 + 11 + 14 + 6 = 44 \quad (3.15)$$

In the case of the profits values of $P2$, the situation is reversed. $A2$ is now the optimal solution and therefore better than $A1$.

The important thing to notice here is that the values for the P matrix right now may appear as having no meaning whatsoever. But we need to think of P as the results obtained from a profit function. Then, we can interpret $P1$ as values of profit in a context where Virgil has just expressed an opinion on social media about the *Divina Commedia* and caused a massive controversy. We can then say that $P1$ is a function which gives more importance to public relations and so the profit p_{21} is very low, whereas $P2$ gives more importance to views and so the profit p_{21} is higher. Of course, in a real problem you know what the function is calculating, but this shows how we can add meaning to a set of symbols in order to understand the data more

efficiently.

With this, we have discussed an assignment problem, looked at its main components and analysed its non-valid, valid and optimal solutions. Some final remarks about the relation between a general assignment problem and the Bootaku problem follow. The actors that perform the jobs, in this case the freelancers that *write* the reviews, are called the *agents*. The *tasks* to be performed are, in the Bootaku problem, the book reviews. Nevertheless, the agents in an assignment problem do not need to be persons (or even things that carry out actions), they can be machines, warehouses, or classrooms. The same can be said for the tasks.

3.2 Greedy algorithms

One way of solving the Bootaku problem described earlier can be found in *greedy algorithms*. A greedy algorithm [GV98] will try to find a subset of candidates that meet the problem constraints and that form the optimal solution. To do so, the algorithm is run iteratively. In each iteration, it will select the best candidate for that precise moment, neglecting future consequences (that is why they are called *greedy*)¹. Before adding a candidate to the solution, the algorithm will determine if it is promising. If the answer is yes, then the candidate is added to the solution. Otherwise, the candidate is no longer evaluated. Each time a candidate is added to the solution, the greedy checks whether the current solution is valid or not.

With this in mind, here follows the *pseudocode* of the generic Greedy Algorithm.

¹For example, let's say I'm walking down the street and I get thirsty. On my mental list of candidate drinks, water has a value of 5 points, lemonade 3 and tea 1. The first vending machine I come across on the street only sells lemonade and tea, so I buy lemonade. However, on the next street I find another vending machine that sells water, but as I am no longer thirsty I don't buy any more drinks. That is, I found a valid solution but not the optimal solution.

Algorithm 1 Generic Greedy Algorithm

```
1: procedure GreedyAlgorithm(candidates)
2:    $x \leftarrow \epsilon$ 
3:    $solution \leftarrow \{\}$ 
4:    $found \leftarrow false$ 
5:   while  $!isEmpty(candidates)$  and  $!found$  do
6:      $x \leftarrow selectCandidate(candidates)$ 
7:     if  $isPromising(x, candidates)$  then
8:        $addToSolution(x, solution)$ 
9:       if  $isSolution(solution)$  then
10:         $found \leftarrow true$ 
11:       end if
12:     end if
13:   end while
14:   return  $solution$ 
15: end procedure
```

As we can see, the generic greedy algorithm has a very simple and elegant definition. However, even though greedy algorithms are easy to implement and can obtain efficient solutions, they are not perfect. Their main flaw relies on their selection function. It is difficult to design a function that can simultaneously find a good local result and translate it into a good global result. That is, the best candidate in some iteration may not be part of the optimal solution.

Next, we will solve the Bootaku problem using a greedy algorithm. The greedy algorithm that we are going to use is inspired by the one defined in [GV98] for solving the *Assignment of tasks* problem.

Algorithm 2 Greedy Algorithm for the Bootaku problem

```
1: procedure GreedyBootaku(profits, assignments)
2:    $best \leftarrow \epsilon$ 
3:   for each freelancer  $F_i \in F$  do
4:     for each review  $R_j \in R$  do
5:        $assignments[F_i, R_j] = false$   $\triangleright$  We initialise the assignments matrix
        (to false or zero, it does not matter).
6:     end for
7:   end for
8:   for each review  $R_j \in R$  do
9:      $best \leftarrow bestFreelancerFor(profits, assignments, R_j)$ 
10:     $assignments[best, R_j] = true$   $\triangleright$  Again, it can be true or one, depending
        on the implementation.
11:   end for
12:   return  $assignments$ 
13: end procedure
```

Algorithm 3 bestFreelancerFor procedure for the Bootaku problem

```
1: procedure bestFreelancerFor(profits, assignments, review)
2:    $best \leftarrow \epsilon$ 
3:    $min \leftarrow$  maximum integer value
4:   for each freelancer  $F_i \in F$  do
5:     if  $profits[F_i, review] < min$  then
6:        $min \leftarrow profits[F_i, review]$ 
7:        $best \leftarrow F_i$ 
8:     end if
9:   end for
10:  return  $best$ 
11: end procedure
```

Those are the two procedures needed in order to solve the Bootaku problem. In the *Assignments of tasks* problem described in the book, the authors define an extra function that checks if the worker is already assigned to another task. However, because our problem is not balanced (we have a different number of tasks and agents), it means that we can have a freelancer writing more than one book review, so that extra function is not required.

The Bootaku problem is really simple to solve because of its lack of constraints. This is done deliberately to focus more on the components of assignment problems and to not waste time on explaining difficult restrictions. However, most problems, including the real problem this document defines (to assign classrooms to the groups of the School), have a lot of constraints.

One way to complicate the Bootaku problem would be to assign completion times to each review. We would have a $n \times m$ T matrix with the completion times for all freelancers and reviews.

$$T = \begin{matrix} & \begin{matrix} DC & EQ & VN & TM \end{matrix} \\ \begin{matrix} Dante \\ Virgil \\ Beatrice \end{matrix} & \begin{pmatrix} t_{11} & t_{12} & t_{13} & t_{14} \\ t_{21} & t_{22} & t_{23} & t_{24} \\ t_{31} & t_{32} & t_{33} & t_{34} \end{pmatrix} \end{matrix} \quad (3.16)$$

Then we could have a maximum time per freelancer. This would force the greedy algorithm to perform a check before assigning a review to a freelancer. If the time it takes to write the review surpasses the maximum time available for that freelancer, the assignment cannot be made.

Let's say that Beatrice has been assigned to the book review for *El Quijote*, so she has already spent a total time of $t_{32} \leq \text{maxTime}$. In a future iteration the greedy algorithm evaluates Beatrice for reviewing *Todo Modo*. Even if she has the greatest profit for *Todo modo*, it is still not enough. The greedy algorithm first has to check in the *bestFreelancerFor* procedure if $t_{32} + t_{34} \leq \text{maxTime}$ and, if the condition is true, then the assignment is performed.

We can notice in the Beatrice example the main failure of greedy algorithms. She was assigned to *El Quijote* for a profit p_{32} and then evaluated again for *Todo modo* with a profit p_{34} . Imagine that she has not enough time left to be able to review the second book and that p_{34} is *way bigger* than p_{32} . This is where assigning the best local result a_{32} would end up ruling out the possibility of assigning the better global result a_{34} .

3.3 Heuristics and metaheuristics

Search strategies in a search problem can be *informed* or *uninformed*. Informed strategies use knowledge specific to a given problem but that is outside of its definition, making this type of strategies more efficient than uninformed strategies. The main way of applying our knowledge of a given problem into the search algorithm designed to solve said problem is by means of *heuristic functions*. An heuristic function $h(n)$ [RN10] represents an estimation of the minimum cost of getting to the objective state from the state given by node n . To expand a node, the algorithm also makes use of an *evaluation function*. An evaluation function $f(n)$ analyses the non-expanded nodes and selects the one with the lowest cost. In the case of greedy algorithms, the evaluation function of a node n is equivalent to the heuristic function of the same node. So we have that $f(n) = h(n)$.

Now that we have explained what heuristic functions are, one question remains. What are *metaheuristics*? Analysing the word, one could think that the *meta* prefix implies that metaheuristics are *heuristics about heuristics*, in the same way *metadata* is *data about data*. However, as Luke [Luk13] points out, this is not the case at all. He defines metaheuristics as:

... a rather unfortunate term often used to describe a major subfield, indeed the primary subfield, of **stochastic optimization**. Stochastic optimization is the general class of algorithms and techniques which employ some degree of randomness to find optimal (or as optimal as possible) solutions to hard problems. Metaheuristics are the most general of these kinds of algorithms, and are applied to a very wide range of problems.

There are many methods of designing algorithms based on *metaheuristics*. In this project we will focus on the Evolutionary Computation method, a subtype of Population-based methods.

3.3.1 Evolutionary Computation

Evolutionary Computation (EC) [Luk13] takes inspiration from population biology, genetics and evolution ². We are interested in the types of algorithms designed using

²Because this method uses vocabulary from these fields of biology, we have followed Luke's approach and defined these terms one by one. A list of definitions of the most commonly used terms in

this method, known as Evolutionary Algorithms (EAs). An EA may be (most of the times) either a *generational algorithm* or a *steady-state algorithm*. A generational algorithm creates a new population of individuals, based on the previous one, in each iteration. Moreover, a steady-state algorithm changes a subset of individuals in each iteration, but not the entire population. The most common EAs are the *Genetic Algorithms* and the *Evolution Strategies*, and there are generational and steady-state versions of the two.

Below is the pseudocode for an abstract generational algorithm.

Algorithm 4 Abstract Generational Algorithm

```

1: procedure AbstractGenerationalAlgorithm(maxTime)
2:    $P \leftarrow$  create the initial population
3:    $best \leftarrow \epsilon$ 
4:    $currentTime \leftarrow$  get current time
5:   while  $\neg idealSolution(P)$  and  $currentTime \leq maxTime$  do
6:      $evaluate(P)$  ▷ Calculate the fitness of all individuals.
7:     for each individual  $P_i \in P$  do
8:       if  $best = \epsilon$  and  $Fitness(P_i) > Fitness(best)$  then
9:          $best \leftarrow P_i$ 
10:      end if
11:    end for
12:     $P \leftarrow newGeneration(P, breed(P))$ 
13:     $currentTime \leftarrow$  update time
14:  end while
15:  return  $best$ 
16: end procedure

```

The initial population in this kinds of algorithms is created by adding random individuals to a set until the maximum population size is reached. Some good practices for this process follow. The most important thing is not generating repeated individuals. This can be done with a dictionary in which we store the individuals as keys. For every new randomly generated individual we check if it is not already contained in said dictionary before adding it to the population set. Finally, it is possible to include individuals designed by hand into the initial population (this is called *seeding* the

EC has been created in the annex of definitions and abbreviations.

population). However, the use of EAs already implies that finding a good heuristic for the problem is not trivial. So even if we think that the individuals we design may be going in the right direction, it is very likely that they will end up producing poor results.

The main difference between generational EAs relies on how they create the new generation. This process is done by means of different operations such as *selection*, *crossover* and *mutation*. Also, some EAs simply discard all the parents in the new generation and others include them again if they have an acceptable fitness. We will take a look at two specific EAs in the following sections, Evolution Strategies and Genetic Algorithms.

3.3.2 Evolution Strategies

Evolution Strategies (ES) [Luk13] are a type of Evolutionary Algorithms. They make use of a selection operator called *Truncation Selection*. This operator consists of selecting individuals from the highest to the lowest fitness value until a predetermined number of selected candidates is reached. For creating the new generation, ES simply use the mutation operator, without combining it with a crossover operator.

The only ES to be covered in this section is the (μ, λ) algorithm. We have chosen it because it is one of the simplest ES and therefore easier to understand. The design of (μ, λ) is essentially one version of the Abstract Generational Algorithm that details the way in which the new generation is built. This new generation is constructed by using both μ and λ parameters.

In this algorithm, we have an initial population of λ number of individuals which are randomly generated. Then, we evaluate the fitness of all individuals, as we did in the Abstract Generational Algorithm, and we calculate the best individual in the generation. The next step is to create the new generation. To do so, (μ, λ) performs the Truncation Selection on the parents, selecting the μ number of parents with greatest fitness, and for each parent a λ/μ number of children are generated. A mutation operation is performed to create the offspring from a copy of their parent. This whole process is done until we arrive at an optimal solution or the maximum time runs out. The pseudocode of the (μ, λ) algorithm is shown as follows.

Algorithm 5 The (μ, λ) Evolution Strategy

```
1: procedure MuLambdaES( $\mu, \lambda$ )
2:    $best \leftarrow \epsilon$ 
3:    $P \leftarrow \{\}$ 
4:   for  $\lambda$  times do
5:      $P \leftarrow P \cup \{\text{new random individual}\}$ 
6:   end for
7:   while  $\neg idealSolution(P)$  and  $currentTime \leq maxTime$  do
8:      $evaluate(P)$  ▷ Calculate the fitness of all individuals.
9:     for each individual  $P_i \in P$  do
10:      if  $best = \epsilon$  and  $Fitness(P_i) > Fitness(best)$  then
11:         $best \leftarrow P_i$ 
12:      end if
13:    end for
14:     $Q \leftarrow \mu$  individuals with the highest to lowest fitness
15:     $P \leftarrow \{\}$ 
16:    for each individual  $Q_j \in Q$  do
17:      for  $\lambda/\mu$  times do
18:         $P \leftarrow P \cup \{mutation(copy(Q_j))\}$ 
19:      end for
20:    end for
21:     $currentTime \leftarrow \text{update time}$ 
22:  end while
23:  return  $best$ 
24: end procedure
```

Knowing how to give values to the parameters λ , μ and the mutation probability is very important in this algorithm. In the case of λ , as it approaches ∞ the algorithm starts behaving as a random search algorithm, so it is best if it does not have an excessively high value. For μ , if we give it a very low value, the algorithm becomes very selective and focuses only of a specific type of individual with a high fitness value. This may result in *premature convergence* of the algorithm and thus end the execution with a locally optimal rather than a globally optimal solution. Lastly, because the mutation operation determines the similarity between parents and offsprings, if the mutation probability is very high, the new population would be very different from the previous one. Therefore, it would make children appear like random individuals.

To conclude the section, simply note that there is another algorithm, very similar to this one, called $(\mu + \lambda)$. The only difference between $(\mu + \lambda)$ and (μ, λ) is that while (μ, λ) discards the parents when creating the new generation, $(\mu + \lambda)$ makes a union between parents and children. This makes each new generation of the $(\mu + \lambda)$ algorithm having a size of $\mu + \lambda$, where μ is the number of parents and λ the number of new offspring. Because very fit parents can survive for several generations, $(\mu + \lambda)$ behaves like a (μ, λ) with a very low μ number, it can terminate with a premature convergence.

3.3.3 Genetic algorithms

The Genetic Algorithm (GA) [Luk13] is a type of Evolutionary Algorithms with a strong similarity towards the (μ, λ) Evolution Strategy. What separates the two algorithms the most is the selection operation and the way a new generation is created. In (μ, λ) , candidates were chosen by Truncated Selection. Once the candidates are selected, the next generation is populated by mutations of the parental copies. However, in the GA, a pair of parents is selected and their offspring are immediately created, adding them to the new generation. This process of simultaneous selection and reproduction occurs until the population reaches the maximum number of individuals set. Further discussion of the GA will follow, once we have seen its pseudocode.

Algorithm 6 The Genetic Algorithm (GA)

```
1: procedure GeneticAlgorithm(popsiz, maxTime)
2:    $best \leftarrow \epsilon$ 
3:    $P \leftarrow \{\}$ 
4:   for  $popsiz$  times do
5:      $P \leftarrow P \cup \{\text{new random individual}\}$ 
6:   end for
7:   while  $\neg idealSolution(P)$  and  $currentTime \leq maxTime$  do
8:      $evaluate(P)$  ▷ Calculate the fitness of all individuals.
9:     for each individual  $P_i \in P$  do
10:      if  $best = \epsilon$  and  $Fitness(P_i) > Fitness(best)$  then
11:         $best \leftarrow P_i$ 
12:      end if
13:    end for
14:     $Q \leftarrow \{\}$  ▷ Here the GA begins to differ from  $(\mu, \lambda)$ .
15:    for  $popsiz/2$  times do
16:      Parent  $P_a \leftarrow selectWithReplacement(P)$ 
17:      Parent  $P_b \leftarrow selectWithReplacement(P)$ 
18:      Children  $C_a, C_b \leftarrow crossover(copy(P_a), copy(P_b))$ 
19:       $Q \leftarrow Q \cup \{mutate(C_a), mutate(C_b)\}$ 
20:    end for
21:     $P \leftarrow Q$ 
22:     $currentTime \leftarrow \text{update time}$ 
23:  end while
24:  return  $best$ 
25: end procedure
```

As we said before, the GA is differentiated from the (μ, λ) ES by means of its selection, crossover and mutation operators. We will now elaborate on these operators, defining them and explaining some of their implementations.

3.3.3.1 Selection

We begin with the selection operator. Even if the selection variant in the GA is different from the one in (μ, λ) , the concept of selection is equivalent in both algorithms. In other words, both use the operation of selection to obtain the parents who will

produce the children of the future generation.. There can be multiple ways to implement the selection operator, including *Random Selection*, *Fitness-Proportionate Selection*, *Stochastic Universal Sampling*, *Tournament Selection* and a variant of the GA which includes *elitism*.

The random selection variant, as its name implies, picks two random parents, removing them from the population. After generating the offspring, another pair of parents is selected and so on, until the new population is complete. In the GA designed to solve the problem of classroom management, we used this variant combined with a tournament between the randomly selected pair of parents and their two generated children to select the best two individuals out of the four (in terms of fitness).

Algorithm 7 Random Selection

```
1: procedure RandomSelection(P)
2:   selected  $\leftarrow$  random individual from population P
3:   remove(selected, P)  $\triangleright$  The individual is removed from the population.
4:   return selected
5: end procedure
```

Next, we will address the topic of Fitness-Proportionate Selection, also known as *Roulette Selection*. In this variant, all individuals are dimensioned according to their fitness. If we think of this process as a lottery, the larger the size of an individual, i.e. the greater the fitness, the more likely the individual is to win the lottery prize. In this case that prize is to be selected to be a parent. This means that a random number n such that $0 \leq n \leq$ the total sum of all fitness values will fall in range of one of the individuals, thus selecting such an individual.

Algorithm 8 Fitness-Proportionate Selection

```
1: procedure GenerationPreparations(P)  ▷ Executed only one time at the start of
   each generation.
2:   global  $\vec{p} \leftarrow \langle p_1, p_2, \dots, p_l \rangle$   ▷ Vector which copies all individuals in P.
3:   global  $\vec{f} \leftarrow \langle f_1, f_2, \dots, p_l \rangle$   ▷ Vector with the fitness values of all the individuals
   in  $\vec{l}$ , keeping the order.
4:   if  $\vec{f}$  contains only zeros then
5:     Substitute all items of  $\vec{f}$  with ones
6:   end if
7:   for  $i$  from 2 to  $l$  do  ▷ Convert  $\vec{f}$  into a cumulative distribution.
8:      $f_i \leftarrow f_i + f_{i-1}$ 
9:   end for
10: end procedure
11: procedure FitnessProportionateSelection(P)
12:    $n \leftarrow$  random number from 0 to  $f_l$  inclusive
13:    $selected \leftarrow p_1$ 
14:   for  $i$  from 2 to  $l$  do
15:     if  $f_{i-1} < n \leq f_i$  then
16:        $selected \leftarrow p_i$ 
17:     end if
18:   end for
19:   return  $selected$ 
20: end procedure
```

We can observe that this operator uses two procedures. The first one is an auxiliary procedure named *GenerationPreparations*, which defines the global vector variables \vec{p} and \vec{f} representing the individuals of the population and their fitness values. The second and main one is the actual selection. In this main procedure we obtain a random number and extract the individual whose range of values contains that chosen random number.

A derivative of the Fitness-Proportionate Selection operator mentioned at the beginning of the section is called Stochastic Universal Sampling (SUS). SUS has two very similar procedures to the ones shown before, one executed one time each generation (normally) that defines some global variables and a main procedure which performs the selection. The main difference comes in how the selection is performed.

Let s be the sum of all fitness values and l be the population size. A random number generated between 0 and s/l will select the individual in that range. For every remaining selection the position value, which started in the random number, is increased by s/l and a new selection is carried out (up to a maximum of l times).

Algorithm 9 Stochastic Universal Sampling Selection

```

1: procedure GenerationPreparations(P)  ▷ Executed only one time at the start of
   each generation.
2:   global  $\vec{p} \leftarrow \langle p_1, p_2, \dots, p_l \rangle$   ▷ Vector which copies all individuals in P.
3:    $\vec{p} \leftarrow shuffle(\vec{p})$ 
4:   global  $\vec{f} \leftarrow \langle f_1, f_2, \dots, p_l \rangle$  ▷ Vector with the fitness values of all the individuals
   in  $\vec{l}$ , keeping the order.
5:   global  $index \leftarrow 0$ 
6:   if  $\vec{f}$  contains only zeros then
7:     Substitute all items of  $\vec{f}$  with ones
8:   end if
9:   for  $i$  from 2 to  $l$  do  ▷ Convert  $\vec{f}$  into a cumulative distribution.
10:     $f_i \leftarrow f_i + f_{i-1}$ 
11:  end for
12:  global  $value \leftarrow$  random number from 0 to  $f_l/l$  inclusive
13: end procedure
14: procedure SUS(P)
15:   while  $f_{index} < value$  do
16:     $index \leftarrow index + 1$ 
17:  end while
18:   $value \leftarrow value + f_l/l$ 
19:  return  $selected$ 
20: end procedure

```

The main advantage of Stochastic Universal Sampling over Fitness-Proportionate selection lies in the fact that while an individual with a high fitness (higher than s/l) might never be chosen in a Fitness-Proportionate selection, in the Stochastic Universal Sampling variant its selection is guaranteed.

The last variant we will explain is the Tournament Selection. In contrast with these past two variants, the Tournament Selection is a very straightforward algorithm. In it, a t number of candidates are selected and the fittest is returned, like a sports

competition. Every time a t_i candidate is selected, it is removed from the population. The pseudocode for this variant is shown below.

Algorithm 10 Tournament Selection

```
1: procedure TournamentSelection( $P, t$ )
2:    $best \leftarrow$  random individual from population  $P$ 
3:    $remove(best, P)$ 
4:   for  $i$  from 2 to  $t$  do
5:      $next \leftarrow$  random individual from population  $P$ 
6:      $remove(next, P)$ 
7:     if  $fitness(next) > fitness(best)$  then
8:        $best \leftarrow next$ 
9:     end if
10:  end for
11:  return  $best$ 
12: end procedure
```

This variant is both simple and flexible. Its flexibility comes from the variable size of candidates for the tournaments. Some considerations for the value of t follow. If t is very low, the operator behaves like a random search. However, if t is very high, the individual with the greatest fitness value will have a much higher likelihood of showing up and getting picked every time. As we stated when talking about Random Selection, a combination of both Random and Tournament selections were implemented in the final design of the GA used in the prototype.

4 Problem definition

The School of Computing Engineering of Oviedo must find a classroom for each group of a given semester. In most cases of this particular problem, just like in the Bootaku situation described in 3, there are more *tasks* (groups/book reviews) than *agents* (classrooms/freelancers). And, in the same way, a valid solution implies that all groups have one (and only one) classroom assigned.

The data for the classrooms and groups can be represented by two sets C and G .

$$C = \{c_1, c_2, \dots, c_n\} \quad (4.1)$$

$$G = \{g_1, g_2, \dots, g_m\} \quad (4.2)$$

Where C is the set of n elements representing all the classrooms of the School, and G a set of m elements representing the groups for a given semester.

A classroom can be a laboratory or a theory classroom. Groups, on the other hand, can be taught in English or Spanish, and have three types. Laboratory, theory and seminar groups. In this problem, because we are only interested in the classrooms that can be assigned to groups, we only consider two types. Laboratory and theory. That is, we consider that the types of classrooms are *identical* to the types of groups.

$$T = \{t_1, t_2, \dots, t_p\} \quad (4.3)$$

$$L = \{l_1, l_2, \dots, l_q\} \quad (4.4)$$

Therefore, each classroom c and group g have a type t . In the case of groups, they also are taught in language l .

Each group has a set of academic weeks and of group schedules. A group can attend classes weekly, every two weeks or on a non-trivial pattern, and may be taught on one or several days.

$$W_i = \{w_{i1}, w_{i2}, \dots, w_{ir}\} \quad (4.5)$$

$$H_i = \{h_{i1}, h_{i2}, \dots, h_{is}\} \quad (4.6)$$

Therefore, every group i has a set of weeks W_i and a set of schedules H_i . A schedule consists of a triplet in the form $(DayOfTheWeek, start(hh : mm), finish(hh : mm))$.

Every group belongs to a subject.

$$S = \{s_1, s_2, \dots, s_t\} \quad (4.7)$$

A subject s is related to a subset of G groups.

With these, we have defined all the data which we need in order to solve the problem. Now we will talk about the problem constraints that we have to fulfill.

For the constraints, we call *hard constraints* those which are imperative for the solution to be valid, and *soft constraints* the ones that reflect on the overall quality of the solution but are not mandatory.

Before introducing the constraints, two new concepts are presented. Restrictions and preferences.

$$R_i = \{r_{i1}, r_{i2}, \dots, r_{iu}\} \quad (4.8)$$

$$P_i = \{p_{i1}, p_{i2}, \dots, p_{iv}\} \quad (4.9)$$

Restrictions and preferences can be positive or negative. A group i must be assigned to a classroom in the set of its positive restrictions, and cannot be assigned to a classroom in the set of its negative restrictions. It is preferred that the group i is assigned to a classroom in the set of its positive preferences, and preferably not in

the set of its negative preferences. With that in mind, the constraints are listed next.

Hard constraints:

Laboratory groups can only be assigned to laboratories.

Theory and seminar groups can only be assigned to theory classrooms.

A group cannot be assigned to a classroom whose capacity is less than the number of students in the group.

A group with a set of positive restrictions must be assigned to one of those classrooms.

A group with a set of negative restrictions cannot be assigned to one of those classrooms.

A group cannot be assigned to a classroom if that classroom was already assigned to another group and both groups collide (they overlap in week and schedule).

Soft constraints:

Laboratory groups of the same subject must all attend the same laboratory classroom, and if not possible, at least minimise the number of laboratories assigned to them.

Theory groups of the same name and course work in the same way, but being assigned to theory classrooms ¹.

English and Spanish groups should go to different classrooms.

Every hour a number of laboratories must be empty. To cover for emergencies.

A group with a set of positive preferences should be assigned to one of those classrooms.

A group with a set of negative preferences should not be assigned to one of those classrooms.

¹All the groups in the School have the format *subject.type.name*. For example the group *Com.T.1* refers to *theory* group *1* of the *Computability* subject. So all theory groups named 1 would be assigned to the same theory classroom, if possible.

5 Proposed solution

5.1 Search space

5.1.1 Assignments

An assignment is a tuple which associates a group with a classroom.

$$(G_i, C_j) \quad (5.1)$$

Because a group can only have *one* classroom assigned, an assignment can be identified by the *code* ¹ of the group. For example, the assignment for group SI.T.1 can be identified by the code SI.T.1 as well.

Assigning just *one* classroom to each group means that the total number of assignments is calculated by the following expression.

$$TotalNumberOfAssignments = |G| \quad (5.2)$$

This implies that there are as many assignments as the number of groups for the semester.

5.1.2 Solutions

A *solution* for this problem is represented by a set of all the assignments must be performed for the semester. As presented in the previous section, the total number of assignments equals the total number of groups in that semester. So we have the next statement.

$$Solution = \{(G_1, C_x), (G_2, C_y), \dots, (G_m, C_z)\} \quad (5.3)$$

Where m is the total number of groups and x, y and z are the index for the class-

¹The name convention previously mentioned: *subject.type.name* (e.g. Com.T.1).

rooms assigned to the groups. Note that the classrooms are not sequential (e.g x could represent C_{12} and y represent C_3).

An *empty solution* is represented by a set of all the assignments where each assignment is *incomplete*. We mean that an assignment is incomplete when the group has no classroom assigned.

$$IncompleteAssignment = (G_i, -) \quad (5.4)$$

So, for the empty solution, we have a set with the following format.

$$EmptySolution = \{(G_1, -), (G_2, -), \dots, (G_m, -)\} \quad (5.5)$$

Finally, a *partial solution* is one in which not every assignment was performed, and a *complete solution* is defined by a set in which all the assignments have been performed and each group has a classroom associated with it.

5.1.3 States

A *state* represents a phase in the problem. There can be three states. The *initial state*, which stands for the empty solution of non allocated assignments. The *final state*, which represents a complete solution with all the assignments performed. And the *intermediate states* portraying partial solutions.

A key concept to understand our solution is the following. Although by default we start the execution of the algorithms with an initial state, it is also possible to start the execution with an intermediate state. This is because we can receive as input a partial or total solution of assignments and work from there. Now we will discuss how we can jump from one state to the next, which is normally called *state expansion*.

To expand a state, one of the non performed assignments in the solution is executed. This means that every time a classroom is assigned to a group the state is being expanded. To perform an assignment, the number of possible candidates is the same as the number of classrooms. So we have the following.

$$TotalNumberOfCandidates = |C| \quad (5.6)$$

Nevertheless, as there exist constraints that indicate whether or not the solution is valid, there are filters which reduce the number of available classrooms for a group. This allows for optimized and easy to retrieve calculations in the execution of the greedy algorithm (this will be explained later in 5.3). The important thing to note at this moment is that, because of these filters, not all states can be expanded.

5.1.4 Instances

The complexity of the calculations and completion time depend on many factors. Some of those factors follow. First, the number of groups for the semester, which directly translates to the number of assignments to be made. Second, the number of classrooms. If there are more classrooms, it is easier to avoid collisions. Obviously, the number of groups is much more volatile between semesters than the number of classes, which is likely to change very occasionally, if at all. Lastly, the case of starting the prototype with an intermediate state. This means that the set of assignments represents a partial solution given as input. The number of calculations decreases in direct proportion to the assignments already made, therefore the completion time would be lower.

Because of the constraints of the problem, there are some groups where class allocations are more straightforward. For example, a group with only one positive restriction is either going to have that classroom assigned to it or, if it collides with other group, end up unallocated. This is why all the groups that just have one available classroom are assigned first. Also, the groups which have less students have more available classrooms than those with a large number of members.

5.2 Collisions

A collision is an overlap of the timetable of two different groups. For a collision to occur, the groups must clash at least once in the same week, day and time. Collisions are an essential part of this problem, as we cannot assign a classroom to a group if another group was previously associated with that classroom and both groups collide.

5.2.1 Lazy Collision Matrix

Due to the large number of assignments that have to be made throughout the execution cycle of the genetic algorithm, the chosen data structures were properly analysed. This is where the *Lazy Collision Matrix* comes in.

Imagine that we have the following group set.

$$G = \{G_1, G_2, G_3\} \quad (5.7)$$

Then, our initial Lazy Collision Matrix would be represented by the expression below.

$$LCM = \begin{matrix} & \begin{matrix} G_1 & G_2 & G_3 \end{matrix} \\ \begin{matrix} G_1 \\ G_2 \\ G_3 \end{matrix} & \begin{pmatrix} & -1 & -1 \\ -1 & & -1 \\ -1 & -1 & \end{pmatrix} \end{matrix} \quad (5.8)$$

First of all, the diagonal is empty because we never compare one group against itself. Then, we can observe that the rest of values are defaulted to -1 . Why? Because there are *not yet evaluated*. That is the reason behind the name of the matrix. It is *lazy* because the collisions are only calculated when needed.

Continuing with this example, imagine that we assign classroom C_x to group G_1 . Then, G_2 also tries to have C_x assigned to it, so we check if both groups collide. We find out that they do, so we update the matrix.

$$LCM = \begin{matrix} & \begin{matrix} G_1 & G_2 & G_3 \end{matrix} \\ \begin{matrix} G_1 \\ G_2 \\ G_3 \end{matrix} & \begin{pmatrix} & 1 & -1 \\ 1 & & -1 \\ -1 & -1 & \end{pmatrix} \end{matrix} \quad (5.9)$$

Therefore, the values are updated with a 1, which indicates that both groups *collide*. This results in a different classroom C_y being assigned to G_2 . Now, G_3 has that classroom also available, so we check if it clashes with G_2 . We learn that they do not

collide. We update the matrix again.

$$LCM = \begin{matrix} & \begin{matrix} G_1 & G_2 & G_3 \end{matrix} \\ \begin{matrix} G_1 \\ G_2 \\ G_3 \end{matrix} & \begin{pmatrix} & 1 & -1 \\ 1 & & 0 \\ -1 & 0 & \end{pmatrix} \end{matrix} \quad (5.10)$$

The value for non-collision is 0, as observed. Because G_3 does not clash with G_2 , they are both allocated in the same classroom.

We can now generalise the LCM as in the next expression.

$$LCM = \begin{matrix} & \begin{matrix} G_1 & G_2 & G_3 \end{matrix} \\ \begin{matrix} G_1 \\ G_2 \\ G_3 \end{matrix} & \begin{pmatrix} & g_{12} & g_{13} \\ g_{21} & & g_{23} \\ g_{31} & g_{32} & \end{pmatrix} \end{matrix} \quad (5.11)$$

Where a value g_{ij} can be

$$\begin{cases} 1, & \text{if } G_i \text{ collides with } G_j \\ 0, & \text{if } G_i \text{ does not collide with } G_j \\ -1, & \text{if the collision has not yet been evaluated} \\ \epsilon, & \text{otherwise} \end{cases}$$

The main advantage of this design is that we do not have to calculate all collisions. For example, a collision between a laboratory group and a theory group would be pointless to calculate because they would never be allocated in the same classroom. Therefore we alleviate the number of calculations.

As there is only one Lazy Collision Matrix, all calculations performed by the greedy algorithm across all populations in all generations are stored in just one place. This means that all collisions are being calculated only when necessary and only once. Think of the previous example. In a future iteration the greedy algorithm wants to check if groups G_1 and G_2 collide. It access the corresponding location in the LCM,

and because it contains a 1, the greedy concludes that they indeed collide. This is done with a $O(1)$ complexity, as the matrix is coded as a dictionary of dictionaries. If instead the greedy wanted to check if groups G_1 and G_3 collided, because the LCM has a -1 in that position, the greedy would have to perform the collision check and then update the matrix.

5.3 Classroom filters

A classroom filter is a function. It receives as input either the C set or a subset $I \subset C$, and outputs a new subset $F \subset C$ with the available classrooms for a given group. It filters out the classrooms that do not comply with the hard constraints for that particular group (except collisions, which are calculated later with the LCM, see 5.2.1). All groups have the same classroom filters.

For example, let us consider a group G_i with type T_j and x students. A type filter for G_i would remove from the result set all the classrooms with a type different from T_j . A capacity filter would then take the set resulting from the previous type filter and use it as input. Then it would eliminate from the set all classrooms with a capacity lower than x and return the new F set.

This reduces the number of classrooms that the greedy has to evaluate and therefore decreases the complexity of the calculations. Furthermore, the filters are deterministic, that is, the execution of the filters of a group will always give the same results. This means that, if needed, the filters are only performed once per group per execution.

This may lead us to this question. *If all classroom filters are executed only once per group every time you run the prototype, why bother using a lazy approach for storing them? Would it not be better to perform and store them in a dictionary at the start of the execution?* The answer is no. It is true that if we start from an empty solution the lazy approach does not present a big advantage. Nonetheless, in the case of partial solutions, it reduces the number of calculations. For example, if we want to assign a classroom to a new group created in the middle of the semester, the only thing we care about is if the group collides with any other group. The filters for the rest of the groups are, in that situation, irrelevant. This is due to the fact that they have already been allocated to their corresponding classrooms.

5.3.1 Lazy Filter Dictionary

The filters work in a similar way to the Lazy Collision Matrix. Again, the election of the data structures is crucial to optimise the execution times. That is why the classroom filters are coded using a dictionary of sets. Once more, we will explain this with an example.

We have a M set of n classroom filters, three groups G_1 , G_2 and G_3 , and two classrooms C_1 and C_2 .

$$M = \{M_1, M_2, \dots, M_n\} \quad (5.12)$$

$$G = \{G_1, G_2, G_3\} \quad (5.13)$$

$$C = \{C_1, C_2\} \quad (5.14)$$

Then, we define the dictionary as a function $Dict : Keys \rightarrow Values \cup \{\epsilon\}$ where ϵ is the *null* character. That is, $\epsilon \notin Values$.

The keys are represented by the groups, so $Keys \equiv G$. The values depict the different sets of available classrooms for each group. Because the dictionary is as lazy as the LCM, the calculations are performed as needed, so the initial set of values are by default ϵ . Then we can say that $Values = \{\epsilon, \epsilon, \epsilon\}$.

Accordingly, we have the following cases.

$$Dict(x) = \begin{cases} \epsilon, & \text{if } x = G_1 \\ \epsilon, & \text{if } x = G_2 \\ \epsilon, & \text{if } x = G_3 \\ \epsilon, & \text{otherwise} \end{cases}$$

We are now in the first execution of the greedy algorithm. Because we start from an empty state and not from an intermediate step, the greedy will try to assign a classroom for all groups. As a result of the values in the dictionary being *null*, the greedy knows that it must execute the filters, updating the dictionary. This is the LFD

after the first execution for this case.

$$Dict(x) = \begin{cases} F_1, & \text{if } x = G_1 \\ F_2, & \text{if } x = G_2 \\ F_3, & \text{if } x = G_3 \\ \epsilon, & \text{otherwise} \end{cases}$$

With each $F_i \subset C$ being the filtered classrooms for each group. Example of values for the F sets follow.

$$F_1 = \{C_1\} \tag{5.15}$$

$$F_2 = \{C_1, C_2\} \tag{5.16}$$

$$F_3 = \{\} \tag{5.17}$$

We can observe that the for G_1 , classroom C_2 was filtered out. In the case of G_2 , both classrooms passed all the constraints. And for G_3 , there are no classes available. This is very important, because it implies that with these filters, G_3 will *always* end up without a classroom. Therefore, a complete solution cannot be found for this case unless the filters are changed.

Lastly, a final remark. We can say that from now on, until the prototype terminates, the greedy will not have to perform the filters for any of the groups again, as the results are already stored in the dictionary.

5.4 Greedy algorithm

5.4.1 Preprocessing

5.4.2 Heuristic

5.4.3 Repairs

5.5 Genetic algorithm

5.5.1 Fitness function

5.5.2 Operators

5.5.2.1 Selection

5.5.2.2 Crossover

5.5.2.3 Mutation

5.5.2.4 Tournament

5.5.3 Parameters

6 Project planning and budget overview

6.1 Planning

6.2 Budget summary

7 Analysis

7.1 System definition

7.2 System requirements

TODO: THESE ARE THE INITIAL REQUIREMENTS, THEY WILL CHANGE

The requirements listed here are a basic overview of the fundamental functionality covered by the project. For the complete list of in-depth requirements the reader might refer to NOPE.

7.2.1 Interface

- The program must implement a CLI.
 - The CLI must show basic or complete information to the user depending on the given option flag.
 - The CLI must show the encountered errors to the user before terminating the execution.
 - The CLI must have help, license and version options.

7.2.2 Input

- The program receives as input the classrooms, groups, group schedule and the academic weeks of each group.
- The program might optionally receive as input a subset of assignments already performed.
- The program might optionally receive as input a previous complete list of assignments but without some of the classrooms/laboratories used in it.
- The program might optionally receive as input a previous complete list of assignments but with more or less groups.
- The program might optionally receive as input a list of classroom preferences for the groups of a particular subject, given their type (theory or laboratory) and language (english or spanish).

7.2.3 Configuration

- Program configuration must allow the user to control the parameters of the genetic algorithm.
- Program configuration must allow the user to change the version of the program.
- Program configuration must allow the user to specify the folder paths for the log and output files.
- Program configuration can change in the middle of the course.

7.2.4 Algorithm

- The program must use a genetic algorithm guided by a greedy algorithm.
- Language group requirements:
 - English groups should go to different classrooms/laboratories from the spanish groups.
- Classroom requirements:
 - Some initial classroom assignments can be specified before the execution of the program and they must remain the same.
 - The program must be able to find a gap in the current list of assignments to include a (mono/multi)-(classroom/laboratory).
 - The number of groups of the same number and course assigned to the same theory classroom must be maximised.
 - The number of groups of the same subject assigned to the same laboratory must be maximised.
 - In each time slot there must be a minimum number of free laboratories.
 - Some big laboratories must be empty for emergency reasons.
 - The program must penalise assignments where the number of students is far below the number of computers.
 - The laboratories must have some free space defined by the user.

- In small laboratories (of 16 computers) there must be at least two free computers.
- The program must be able to handle a split in two of a laboratory group with only one professor (for emergencies).

7.3 Subsystem mapping

7.4 Preliminary class diagram

7.5 Analysis of use cases

7.6 Analysis of user interfaces

7.7 Test plan specification

8 System design

8.1 System architecture

8.2 Class design

8.3 Interaction and state diagrams

8.4 Activity diagram

8.5 Interface design

8.6 Technical specification of the test plan

9 System implementation

9.1 Standards and references

9.1.1 Standards

9.1.2 Licenses

The software of this project is licensed under the GNU General Public License v2.0.

9.1.3 Other references

Java Code Conventions. Set of guidelines and conventions for programmers to consider when using the Java programming language.

9.2 Programming languages

There were two programming languages considered for the implementation of the system, **C** and **Java**.

Considering that:

- The author and only developer of the system has worked with Java throughout his university studies, but only used C in one subject and in some of his personal projects.
- Java is probably less efficient than C when executing the genetic and greedy algorithms.
- Java code is more easy to run in other systems than C code.
- The program is going to be executed only a few times a year.

For this reasons, even if C would be faster in execution, because the program will not be running every day, and taking into account the other two advantages, Java was the language of choice for implementing the system.

9.3 Tools and programs used in development

9.4 System development

10 Test development

10.1 Unit tests

10.2 Integration and system tests

10.3 Usability and accessibility tests

10.4 Performance tests

11 Experimental results

12 System manuals

12.1 Installation manual

12.2 Execution manual

12.3 User manual

12.4 Programmer manual

13 Conclusions and future work

13.1 Final conclusions

13.2 Future work

14 Budget

14.1 Internal budget

14.2 Client budget

15 Annexes

15.1 Definitions and abbreviations

Listed below is a glossary of definitions and abbreviations used in the document whose meaning may not be obvious.

Glossary of definitions:

- **Evolutionary Computation:** method of designing a metaheuristic algorithm. It is a subtype of Population-based methods. The definitions for the common components of evolutionary computation follow [Luk13].
 - **Breeding:** the act of creating one or more children from a population of parents by combining the crossover and mutation operators.
 - **Chromosome:** a specific type of genome consisting of a fixed-length array.
 - **Child and parent:** both are individuals. A child being a possible modification of its parent.
 - **Crossover:** operator that creates childs from parents by means of combining sections of the genomes of the parents.
 - **Evaluation:** calculating the fitness of an individual.
 - **Fitness:** quality of an individual.
 - **Generation:** the population of a given iteration of the algorithm. The next generation is created by means of the different operations defined by said algorithm.
 - **Genome:** the data structure that defines an individual.
 - **Individual:** candidate solution for the problem.
 - **Mutation:** operator that modifies the genome of an individual.
 - **Population:** set of individuals.
 - **Selection:** operator that elects individuals from the population based on some criteria.
- **Genetic algorithm:** metaheuristic search and optimization algorithm.

- **Greedy algorithm:** algorithm that builds the solution in successive steps, always trying to take the optimal solution for each step
- **Heuristic:** function that gives value to each path from a intermediate state to the goal state. Applied in search algorithms, heuristics are based on knowledge outside the problem definition.
- **Java:** general-purpose, high-level, object-oriented programming language.
- **Metaheuristic:** algorithm that uses randomness to find a possible optimal solution to a hard problem. They are part of the stochastic optimization field.

Glossary of abbreviations:

- **CSV:** Comma-Separated Values. Refers to a text file format.
- **CLI:** Command Line Interface.
- **EA:** Evolutionary Algorithm.
- **EC:** Evolutionary Computation.
- **ES:** Evolution Strategies.
- **GA:** Genetic Algorithm.
- **LCM:** Lazy Collision Matrix (see [5.2.1](#)).
- **LFD:** Lazy Filter Dictionary (see [5.3.1](#)).
- **TXT:** Text. Refers to the text file format.

15.2 Submission contents

16 Source code

Bibliography

- [dlC18] Gonzalo de la Cruz. Metaheuristics for the assignment of students to class groups. End-of-degree thesis, School of Computing Engineering of Oviedo, 2018.
- [GV98] Rosa Guerequeta and Antonio Vallecillo. *Técnicas de Diseño de Algoritmos*. Servicio de publicaciones de la Universidad de Málaga, 1998.
- [Luk13] Sean Luke. *Essentials of Metaheuristics*. Lulu, second edition, 2013. Available for free at <http://cs.gmu.edu/~sean/book/metaheuristics/>.
- [Red] Jose Manuel Redondo. *Documentos-modelo para Trabajos de Fin de Grado/Master de la Escuela de Informática de Oviedo*. Escuela Universitaria de Ingeniería Técnica en Informática de Oviedo, 1.4th edition.
- [RN10] Stuart J. Russell and Peter Norvig. *Artificial Intelligence: A Modern Approach*. Pearson Education, third edition, 2010.